

In []:

In []:

1. Now, let's call the particular libraries :

page 74 book 4

```
In [228]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
```

2. Now Read the DataSet :

page 74 book 4

```
In [61]: df = pd.read_csv("DataS/USA_Housing.csv")
df
```

Out[61]:

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price	Address
0	79545.458574	5.682861	7.009188	4.09	23086.800503	1.059034e+06	208 Michael Ferry Apt. 674\nLaurabury, NE 3701...
1	79248.642455	6.002900	6.730821	3.09	40173.072174	1.505891e+06	188 Johnson Views Suite 079\nLake Kathleen, CA...
2	61287.067179	5.865890	8.512727	5.13	36882.159400	1.058988e+06	9127 Elizabeth Stravenue\nDanieltown, WI 06482...
3	63345.240046	7.188236	5.586729	3.26	34310.242831	1.260617e+06	USS Barnett\nFPO AP 44820
4	59982.197226	5.040555	7.839388	4.23	26354.109472	6.309435e+05	USNS Raymond\nFPO AE 09386
...
4995	60567.944140	7.830362	6.137356	3.46	22837.361035	1.060194e+06	USNS Williams\nFPO AP 30153-7653
4996	78491.275435	6.999135	6.576763	4.02	25616.115489	1.482618e+06	PSC 9258, Box 8489\nAPO AA 42991-3352
4997	63390.686886	7.250591	4.805081	2.13	33266.145490	1.030730e+06	4215 Tracy Garden Suite 076\nJoshualand, VA 01...
4998	68001.331235	5.534388	7.130144	5.44	42625.620156	1.198657e+06	USS Wallace\nFPO AE 73316
4999	65510.581804	5.992305	6.792336	4.07	46501.283803	1.298950e+06	37778 George Ridges Apt. 509\nEast Holly, NV 2...

5000 rows × 7 columns

3. Now, TAKING 'TOP 5' AND 'TAIL 5' RECORDS :

PAGE 75 BOOK 4

In [62]: `df.head()`

Out[62]:

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price	Address
0	79545.458574	5.682861	7.009188	4.09	23086.800503	1.059034e+06	208 Michael Ferry Apt. 674\nLaurabury, NE 3701...
1	79248.642455	6.002900	6.730821	3.09	40173.072174	1.505891e+06	188 Johnson Views Suite 079\nLake Kathleen, CA...
2	61287.067179	5.865890	8.512727	5.13	36882.159400	1.058988e+06	9127 Elizabeth Stravenue\nDanieltown, WI 06482...
3	63345.240046	7.188236	5.586729	3.26	34310.242831	1.260617e+06	USS Barnett\nFPO AP 44820
4	59982.197226	5.040555	7.839388	4.23	26354.109472	6.309435e+05	USNS Raymond\nFPO AE 09386

In [63]: `df.tail()`

Out[63]:

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price	Address
4995	60567.944140	7.830362	6.137356	3.46	22837.361035	1.060194e+06	USNS Williams\nFPO AP 30153-7653
4996	78491.275435	6.999135	6.576763	4.02	25616.115489	1.482618e+06	PSC 9258, Box 8489\nAPO AA 42991-3352
4997	63390.686886	7.250591	4.805081	2.13	33266.145490	1.030730e+06	4215 Tracy Garden Suite 076\nJoshualand, VA 01...
4998	68001.331235	5.534388	7.130144	5.44	42625.620156	1.198657e+06	USS Wallace\nFPO AE 73316
4999	65510.581804	5.992305	6.792336	4.07	46501.283803	1.298950e+06	37778 George Ridges Apt. 509\nEast Holly, NV 2...

4. If we want only '1' record :

```
In [64]: df.head(1)
```

```
Out[64]:
```

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price	Address
0	79545.458574	5.682861	7.009188	4.09	23086.800503	1.059034e+06	208 Michael Ferry Apt. 674\nLaurabury, NE 3701...

```
In [65]: df.tail(1)
```

```
Out[65]:
```

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price	Address
4999	65510.581804	5.992305	6.792336	4.07	46501.283803	1.298950e+06	37778 George Ridges Apt. 509\nEast Holly, NV 2...

5. For Information : info()

page 75 book 4

```
In [66]: # Here, Range Index : 5000 entries, (total 7 columns),  
# and There are 'no nulls', Data is Fully Perfect.
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 5000 entries, 0 to 4999
```

```
Data columns (total 7 columns):
```

#	Column	Non-Null Count	Dtype
0	Avg. Area Income	5000 non-null	float64
1	Avg. Area House Age	5000 non-null	float64
2	Avg. Area Number of Rooms	5000 non-null	float64
3	Avg. Area Number of Bedrooms	5000 non-null	float64
4	Area Population	5000 non-null	float64
5	Price	5000 non-null	float64
6	Address	5000 non-null	object

```
dtypes: float64(6), object(1)
```

```
memory usage: 273.6+ KB
```

6. Now import 'Seaborn' :

page 75 book 4

```
In [67]: import seaborn as sns  
sns.heatmap(df.corr(), annot = True)
```

```
Out[67]: <AxesSubplot:>
```



7. Column Names :

page 76 book 4

```
In [68]: # Now i want to view the column names :
```

```
In [69]: df.columns
```

```
Out[69]: Index(['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of Rooms',  
              'Avg. Area Number of Bedrooms', 'Area Population', 'Price', 'Address'],  
              dtype='object')
```

8. Split the Data 'X' 'y' :

page 76 book 4

```
In [71]: X = df[['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of Rooms',  
               'Avg. Area Number of Bedrooms', 'Area Population']]  
y = df['Price']
```

9. 'sklearn library' == SCIENTIFIC LIBRARY :

PAGE 77 BOOK 4

Split X,y into 'train data' and 'test data' :

```
In [78]: # (from sklearn.cross_validation import train_test_split) - This declaration not working :
```

```
In [72]: from sklearn.model_selection import train_test_split
```

train_test_split(|)


```
In [81]: # Here,
```

```
In [82]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_state=42)
```

10. Calling LINEAR REGRESSION :

PAGE 78 BOOK 4

```
In [85]: # Here, Directly Calling Linear Regression :  
# L & R are Capital :  
  
from sklearn.linear_model import LinearRegression
```

```
In [86]: lm = LinearRegression()
```

```
In [87]: lm.fit(X_train,y_train)
```

```
Out[87]: LinearRegression()
```

```
In [88]: print(lm.intercept_)
```

-2638142.110428682

```
In [89]: lm.coef_
```

```
Out[89]: array([2.15898874e+01, 1.66102501e+05, 1.19895936e+05, 1.90107101e+03,  
               1.52315025e+01])
```

```
In [90]: df.columns
```

```
Out[90]: Index(['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of Rooms',  
              'Avg. Area Number of Bedrooms', 'Area Population', 'Price', 'Address'],  
              dtype='object')
```

```
In [92]: cdf = pd.DataFrame(lm.coef_,X.columns, columns = ['Coeff'])
```

```
In [94]: # Here, Based on the input data, The Output data is affected . All 'X' Variables are Independent.
```

```
cdf
```

```
Out[94]:
```

	Coeff
Avg. Area Income	21.589887
Avg. Area House Age	166102.501246
Avg. Area Number of Rooms	119895.936402
Avg. Area Number of Bedrooms	1901.071012
Area Population	15.231503

11. How to Call the Data Set :

page 81 book 4

```
In [ ]: # NOTE :
```

```
# 'boston' - Boston DataSet is an an a 'Inbuilt Data Set', Which is in the form of an a 'Dictionary'.  
# first we need to make ''
```

```
In [96]: from sklearn.datasets import load_boston
```

```
In [99]: df = load_boston()
```

```
In [101]: # Here, 'data'- is an 'X' data, and 'target' - is an a 'y' data.  
# 'feature_names' - are "Labels of 'X' and 'y' ".  
# Here, Data is in a Dictionary Pattern. Now this is Perfect Data Set :  
  
df.keys()
```

```
Out[101]: dict_keys(['data', 'target', 'feature_names', 'DESCR', 'filename', 'data_module'])
```

12. Linear Regression applying on "Sales based on - TV Marketing Budget" :

page 83 book 4

```
In [132]: import pandas as pd  
import numpy as np
```

```
In [141]: advertising = pd.read_csv("DataS/tvmarketing.csv")
advertising
```

Out[141]:

	TV	Sales
0	230.1	22.1
1	44.5	10.4
2	17.2	9.3
3	151.5	18.5
4	180.8	12.9
...
195	38.2	7.6
196	94.2	9.7
197	177.0	12.8
198	283.6	25.5
199	232.1	13.4

200 rows × 2 columns

```
In [142]: advertising.head()
```

Out[142]:

	TV	Sales
0	230.1	22.1
1	44.5	10.4
2	17.2	9.3
3	151.5	18.5
4	180.8	12.9

```
In [143]: advertising.tail()
```

```
Out[143]:
```

	TV	Sales
195	38.2	7.6
196	94.2	9.7
197	177.0	12.8
198	283.6	25.5
199	232.1	13.4

```
In [144]: # Here, We don't have any Nulls(NaN's) : So, the Data is Very Perfect :
```

```
advertising.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0    TV      200 non-null         float64
1   Sales   200 non-null         float64
dtypes: float64(2)
memory usage: 3.2 KB
```

```
In [145]: advertising.columns
```

```
Out[145]: Index(['TV', 'Sales'], dtype='object')
```

describe():

page 84 book 4

```
In [146]: advertising.describe()
```

Out[146]:

	TV	Sales
count	200.000000	200.000000
mean	147.042500	14.022500
std	85.854236	5.217457
min	0.700000	1.600000
25%	74.375000	10.375000
50%	149.750000	12.900000
75%	218.825000	17.400000
max	296.400000	27.000000

importing libraries :

```
In [211]: import seaborn as sns
           %matplotlib inline
           import matplotlib.pyplot as plt
```


In [212]: *# In this particular example, we are going to observe the data and their relation,
by using Seaborn pairlot :*

```
sns.pairplot(advertising,x_vars = ['TV'], y_vars = 'Sales',size = 7,aspect = 0.7, kind = 'scatter')
```

C:\Users\my pc\anaconda3\lib\site-packages\seaborn\axisgrid.py:2076: UserWarning: The `size` parameter has been renamed to `height`; please update your code.

warnings.warn(msg, UserWarning)

Out[212]: <seaborn.axisgrid.PairGrid at 0x2402e1a8e20>



Now, Split Only 'TV' Data :

page 86 book 4

```
In [213]: # Capital 'X' :
```

```
X = advertising['TV']  
X.head()
```

```
Out[213]: 0    230.1  
         1     44.5  
         2     17.2  
         3    151.5  
         4    180.8  
         Name: TV, dtype: float64
```

'Sales Data' :

```
In [214]: y = advertising['Sales']  
y.head()
```

```
Out[214]: 0     22.1  
         1     10.4  
         2      9.3  
         3     18.5  
         4     12.9  
         Name: Sales, dtype: float64
```

Now Train the Model :

```
In [270]: from sklearn.model_selection import train_test_split
```

```
In [271]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_state=42)
```

Types of an a 'X' and 'y' : 'PANDAS' CORE SERIES DATA :

```
In [272]: print(type(X_train))
print(type(y_train))
print(type(X_test))
print(type(y_test))
```

```
<class 'pandas.core.series.Series'>
<class 'pandas.core.series.Series'>
<class 'pandas.core.series.Series'>
<class 'pandas.core.series.Series'>
```

'SHAPE' of an a 'X' and 'y' : Here it is '1' Dimension Data :

PAGE 88 BOOK 4

```
In [259]: # Here, It is '1' Dimensional Data :
```

```
print(X_train.shape)
print(y_train.shape)
print(X_test.shape)
print(y_test.shape)
```

```
(134,)
(134,)
(66,)
(66,)
```

Now, Changing X,y into '2' Dimensional Data :

By using [: , np.newaxis]

page 89 book 4

```
In [273]: X_train = X_train[:,np.newaxis]
X_test = X_test[:,np.newaxis]
```

C:\Users\my pc\AppData\Local\Temp\ipykernel_20152\2733214152.py:1: FutureWarning: Support for multi-dimensional indexing (e.g. `obj[:, None]`) is deprecated and will be removed in a future version. Convert to a numpy array before indexing instead.

```
X_train = X_train[:,np.newaxis]
```

C:\Users\my pc\AppData\Local\Temp\ipykernel_20152\2733214152.py:2: FutureWarning: Support for multi-dimensional indexing (e.g. `obj[:, None]`) is deprecated and will be removed in a future version. Convert to a numpy array before indexing instead.

```
X_test = X_test[:,np.newaxis]
```

```
In [274]: # Here, We are Converting '1' Dimensional to '2' Dimensional Data :
```

```
print(X_train.shape)
print(y_train.shape)
print(X_test.shape)
print(y_test.shape)
```

```
(134, 1)
(134,)
(66, 1)
(66,)
```

Calling 'Linear Regression' :

page 90 book 4

```
In [275]: from sklearn.linear_model import LinearRegression
```

```
In [276]: # 'lr' - Some Left hand Variable :
```

```
lr = LinearRegression()
```

```
In [277]: lr.fit(X_train,y_train)
```

```
Out[277]: LinearRegression()
```

```
In [278]: print(lr.intercept_)
```

```
7.066582521696442
```

'Coefficient' :

```
In [279]: print(lr.coef_)
```

```
[0.04822451]
```

'Prediction' :

```
In [280]: x_pred = lr.predict(X_test)
x_pred
```


```
Out[280]: array([14.94164546, 16.48965233, 21.19154233,  7.63080932, 17.69044269,
 10.68824343, 17.5216569 ,  9.47780816, 17.79171417, 15.51069472,
  8.58565468,  9.77680014, 18.37523078,  7.32699489, 13.79390206,
 15.2743946 ,  7.41862147, 16.59574625, 10.69788834, 18.51508186,
 18.13410821, 10.33620449,  8.89911401, 19.16611279, 10.39407391,
  9.64659396, 17.36251601, 13.78425716, 11.27176004,  7.47166843,
 16.70184018, 10.40371881, 16.65843812,  7.88157679, 20.60320327,
 18.55366147,  9.39582649, 20.26563168, 12.71849542,  8.39275662,
 12.66062601, 15.60232129,  9.14505902, 10.07096967, 17.86887339,
  8.90875891, 10.47123313, 14.17005326, 12.11086656, 10.74611285,
 10.83773942, 15.18759048,  7.48613578,  7.44273372, 10.7509353 ,
 13.30683448, 10.60626176, 21.03722388,  8.01178297, 16.59574625,
 20.77681152, 15.98329494, 12.51113002,  8.21432593, 21.0854484 ,
  8.00213807])
```

ACTUAL DATA VS PREDICT DATA BY 'GRAPH' :

PAGE 91 BOOK 4

```
In [281]: import matplotlib.pyplot as plt
%matplotlib inline
y_pred = lr.predict(X_test)
c = [ i for i in range(1,67,1)]
fig = plt.figure()
plt.plot(c, y_test,color = 'blue', lw = 2.5, linestyle = '-')
plt.plot(c, y_pred,color = 'r', lw = 2.5, linestyle = '-')
```

Out[281]: [<matplotlib.lines.Line2D at 0x2402bca6040>]



NOW, 'KC HOUSING DATA' :

In [566]: *# Once we get this, We can do an a Graphical Representation :*

```
data = pd.read_csv("DataS/29.kc_house_data.csv")
```

In [539]: data.head()

Out[539]:

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	...	grade	sqft_above	sqft_base
0	7129300520	20141013T000000	221900.0	3	1.00	1180	5650	1.0	0	0	...	7	1180	
1	6414100192	20141209T000000	538000.0	3	2.25	2570	7242	2.0	0	0	...	7	2170	
2	5631500400	20150225T000000	180000.0	2	1.00	770	10000	1.0	0	0	...	6	770	
3	2487200875	20141209T000000	604000.0	4	3.00	1960	5000	1.0	0	0	...	7	1050	
4	1954400510	20150218T000000	510000.0	3	2.00	1680	8080	1.0	0	0	...	8	1680	

5 rows × 21 columns



```
In [540]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21613 entries, 0 to 21612
Data columns (total 21 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   id                    21613 non-null  int64  
 1   date                  21613 non-null  object  
 2   price                 21613 non-null  float64 
 3   bedrooms              21613 non-null  int64  
 4   bathrooms             21613 non-null  float64 
 5   sqft_living           21613 non-null  int64  
 6   sqft_lot              21613 non-null  int64  
 7   floors                21613 non-null  float64 
 8   waterfront            21613 non-null  int64  
 9   view                  21613 non-null  int64  
10   condition             21613 non-null  int64  
11   grade                 21613 non-null  int64  
12   sqft_above            21613 non-null  int64  
13   sqft_basement         21613 non-null  int64  
14   yr_built              21613 non-null  int64  
15   yr_renovated          21613 non-null  int64  
16   zipcode               21613 non-null  int64  
17   lat                   21613 non-null  float64 
18   long                  21613 non-null  float64 
19   sqft_living15         21613 non-null  int64  
20   sqft_lot15            21613 non-null  int64  
dtypes: float64(5), int64(15), object(1)
memory usage: 3.5+ MB
```

```
In [541]: data.columns
```

```
Out[541]: Index(['id', 'date', 'price', 'bedrooms', 'bathrooms', 'sqft_living',
                'sqft_lot', 'floors', 'waterfront', 'view', 'condition', 'grade',
                'sqft_above', 'sqft_basement', 'yr_built', 'yr_renovated', 'zipcode',
                'lat', 'long', 'sqft_living15', 'sqft_lot15'],
                dtype='object')
```


In [542]: `data.describe()`

Out[542]:

	id	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	cond
count	2.161300e+04	2.161300e+04	21613.000000	21613.000000	21613.000000	2.161300e+04	21613.000000	21613.000000	21613.000000	21613.000000
mean	4.580302e+09	5.400881e+05	3.370842	2.114757	2079.899736	1.510697e+04	1.494309	0.007542	0.234303	3.400000
std	2.876566e+09	3.671272e+05	0.930062	0.770163	918.440897	4.142051e+04	0.539989	0.086517	0.766318	0.650000
min	1.000102e+06	7.500000e+04	0.000000	0.000000	290.000000	5.200000e+02	1.000000	0.000000	0.000000	1.000000
25%	2.123049e+09	3.219500e+05	3.000000	1.750000	1427.000000	5.040000e+03	1.000000	0.000000	0.000000	3.000000
50%	3.904930e+09	4.500000e+05	3.000000	2.250000	1910.000000	7.618000e+03	1.500000	0.000000	0.000000	3.000000
75%	7.308900e+09	6.450000e+05	4.000000	2.500000	2550.000000	1.068800e+04	2.000000	0.000000	0.000000	4.000000
max	9.900000e+09	7.700000e+06	33.000000	8.000000	13540.000000	1.651359e+06	3.500000	1.000000	4.000000	5.000000

In [543]: `import seaborn as sns`
`import matplotlib.pyplot as plt`
`%matplotlib inline`

In [544]: `from sklearn.model_selection import train_test_split`

In [545]: `X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_state=42)`

In [560]: data

Out[560]:

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	...	grade	sqft_above	sqft_l
0	7129300520	20141013T000000	221900.0	3	1.00	1180	5650	1.0	0	0	...	7	1180	
1	6414100192	20141209T000000	538000.0	3	2.25	2570	7242	2.0	0	0	...	7	2170	
2	5631500400	20150225T000000	180000.0	2	1.00	770	10000	1.0	0	0	...	6	770	
3	2487200875	20141209T000000	604000.0	4	3.00	1960	5000	1.0	0	0	...	7	1050	
4	1954400510	20150218T000000	510000.0	3	2.00	1680	8080	1.0	0	0	...	8	1680	
...
21608	263000018	20140521T000000	360000.0	3	2.50	1530	1131	3.0	0	0	...	8	1530	
21609	6600060120	20150223T000000	400000.0	4	2.50	2310	5813	2.0	0	0	...	8	2310	
21610	1523300141	20140623T000000	402101.0	2	0.75	1020	1350	2.0	0	0	...	7	1020	
21611	291310100	20150116T000000	400000.0	3	2.50	1600	2388	2.0	0	0	...	8	1600	
21612	1523300157	20141015T000000	325000.0	2	0.75	1020	1076	2.0	0	0	...	7	1020	

21613 rows × 21 columns



```
In [547]: import pandas as pd
import numpy as np
from numpy.random import randn
```

```
In [548]: from plotly import __version__
```


```
In [549]: import cufflinks as cf
```

```
In [550]: init_notebook_mode(connected = True)
```


```
-----  
NameError                                Traceback (most recent call last)  
~\AppData\Local\Temp\ipykernel_20152\1831628520.py in <module>  
----> 1 init_notebook_mode(connected = True)  
  
NameError: name 'init_notebook_mode' is not defined
```

```
In [551]: cf.go_offline()
```

```
In [556]: data.iplot()
```

[Export to plot.ly »](#)

```
In [561]: data.iplot(kind = 'scatter')
```



```
In [562]: data = pd.DataFrame({'category': ['A', 'B', 'C'], 'values': [32, 43, 50]})
data
```

Out[562]:


	category	values
0	A	32
1	B	43
2	C	50

```
In [563]: data = pd.DataFrame(np.random.randn(200,4), columns = 'A B C D'.split())  
data.head()
```

Out[563]:


	A	B	C	D
0	1.235486	-0.853671	-0.565666	1.354380
1	0.937049	2.451492	0.738015	1.115578
2	-0.012768	1.446495	1.210544	0.554312
3	-0.001110	-1.967396	1.091417	0.921963
4	0.467174	-0.432148	1.741814	-0.413566

```
In [564]: data.heatmap(kind = 'surface')
```



```
In [567]: import seaborn as sns
sns.heatmap(data.corr(), annot = True)
```

Out[567]: <AxesSubplot:>




```
In [568]: data = pd.DataFrame(np.random.randn(200,4), columns = 'A B C D'.split())  
data.head()
```

Out[568]:

	A	B	C	D
0	-0.717910	0.661480	1.675569	0.541954
1	0.512613	-1.774907	0.839159	-0.899096
2	0.769996	0.052931	-0.901385	-0.510669
3	-0.214891	0.473436	0.169547	0.502355
4	-0.496329	0.223709	-1.580703	-1.568024


```
In [569]: import seaborn as sns  
sns.heatmap(data.corr(), annot = True)
```

Out[569]: <AxesSubplot:>



```
In [570]: data = pd.DataFrame(np.random.randn(200,4))  
data
```


Out[570]:

	0	1	2	3
0	0.909625	1.938965	-0.386127	-1.689274
1	-0.521288	-1.218318	1.512162	2.396411
2	-0.119257	0.716245	-0.542812	-0.251244
3	-0.756151	-0.059014	0.193390	0.002333
4	-0.020800	2.585580	-1.438085	-2.467507
...
195	-0.134047	-2.277579	2.740054	0.063748
196	-0.821618	-1.169159	-0.585291	-1.503375
197	1.448589	1.940807	1.409744	-0.436775
198	0.064037	-1.050196	1.764843	-0.000398
199	-0.381303	-1.611617	-0.092121	-0.153335

200 rows × 4 columns

```
In [571]: import seaborn as sns  
sns.heatmap(data.corr(), annot = True)
```

Out[571]: <AxesSubplot:>



In [572]: *# just for Example :*

```
data = pd.DataFrame(np.random.randn(200,4), columns = 'id price bedrooms bathrooms'.split())  
data.head()
```

Out[572]:

	id	price	bedrooms	bathrooms
0	1.700618	-0.148633	-0.480701	0.208101
1	0.523194	-0.912204	0.267056	0.278257
2	-0.070691	0.184255	0.780837	-0.362575
3	-0.053029	-0.302490	-0.339402	-1.519880
4	1.410595	0.631424	-1.381332	-0.213058

```
In [573]: import seaborn as sns  
sns.heatmap(data.corr(), annot = True)
```

Out[573]: <AxesSubplot:>

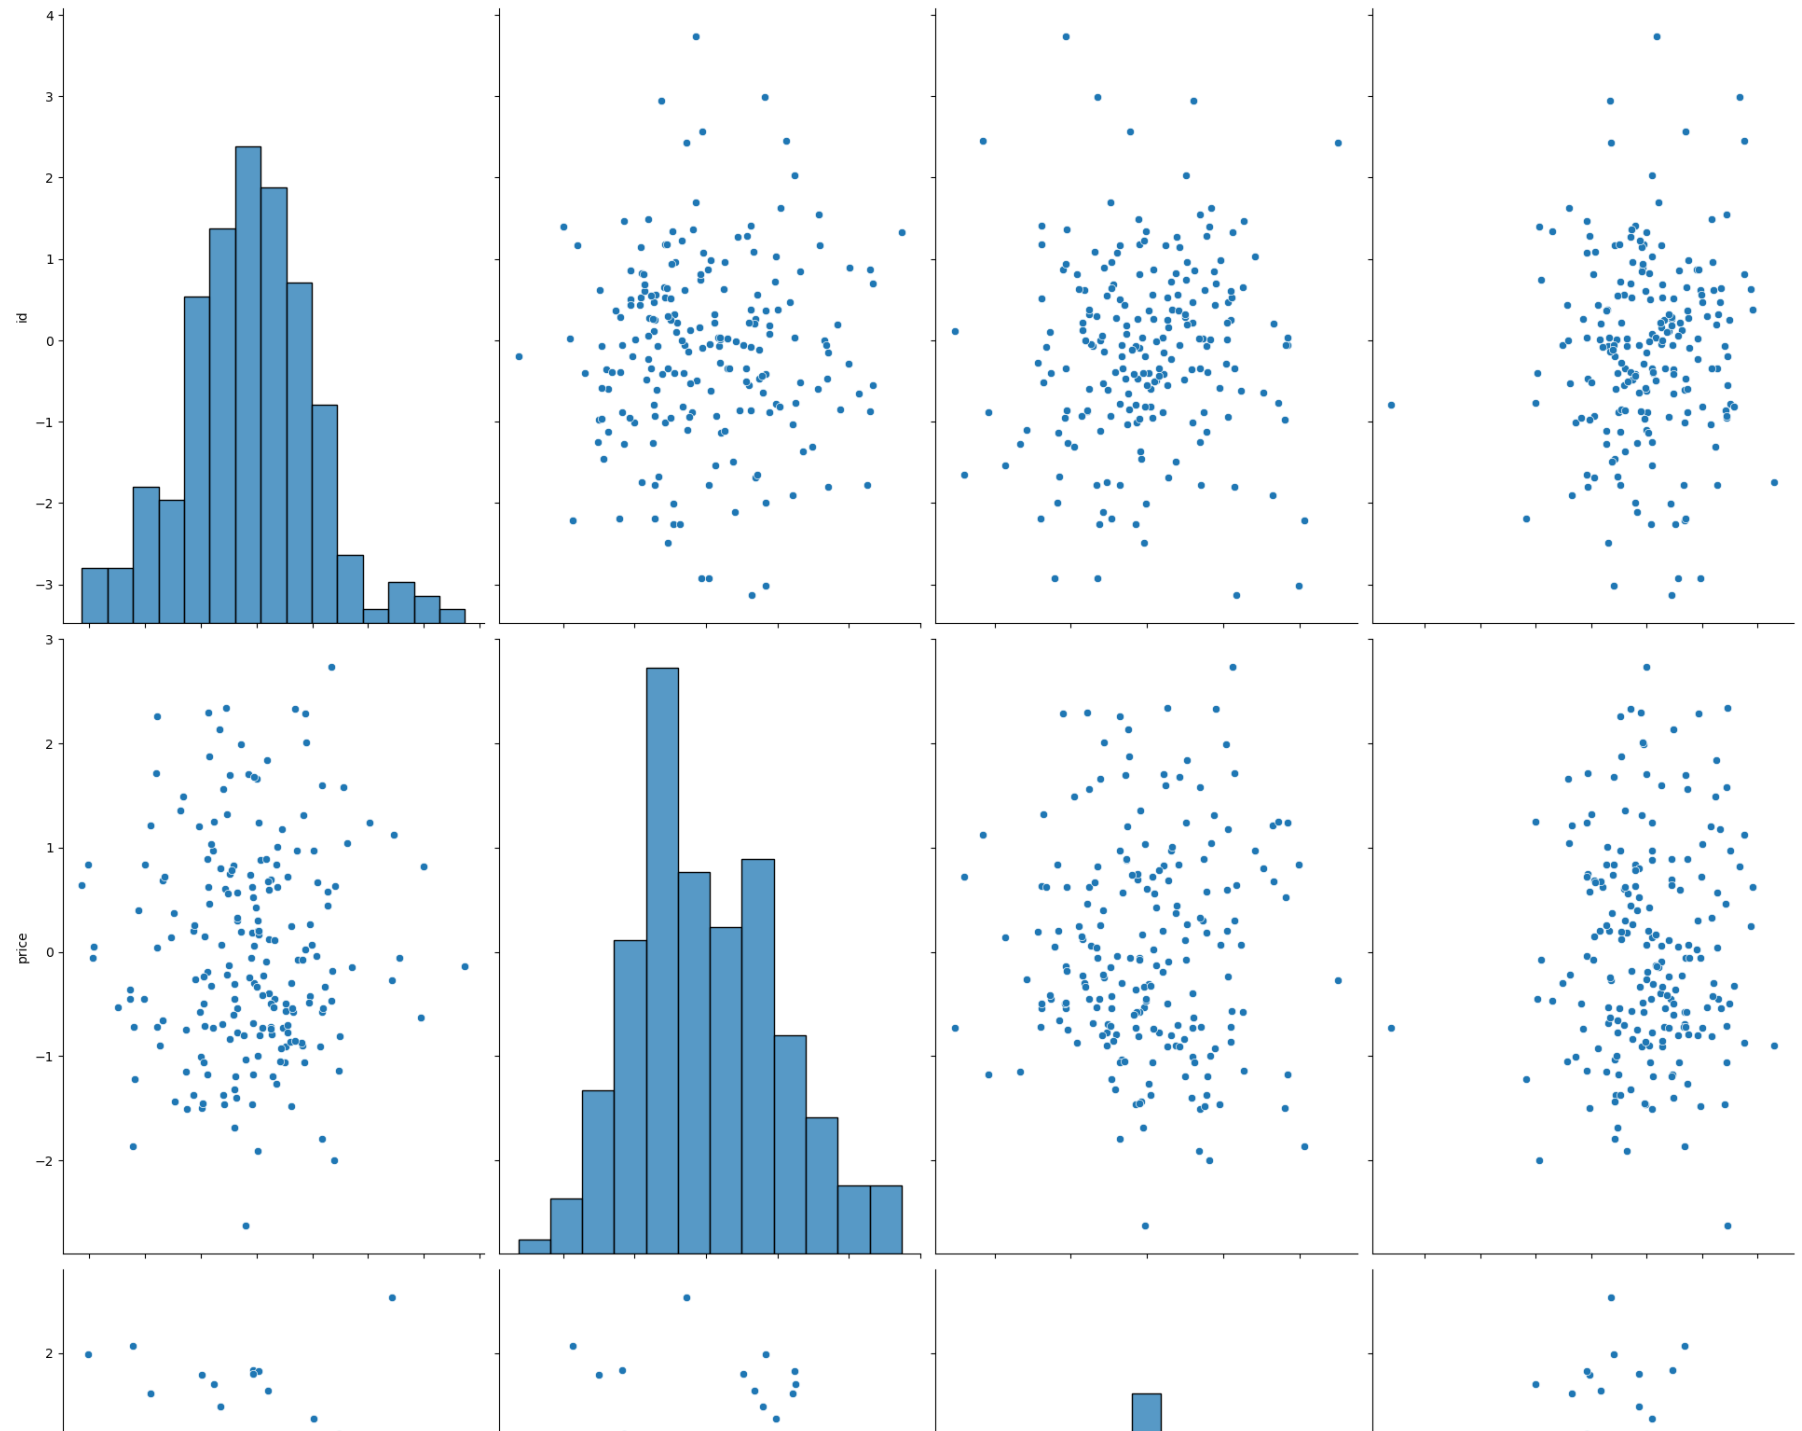


```
In [574]: sns.pairplot(data, size = 7, aspect = 0.7, kind = 'scatter')
```

C:\Users\my pc\anaconda3\lib\site-packages\seaborn\axisgrid.py:2076: UserWarning:

The `size` parameter has been renamed to `height`; please update your code.

```
Out[574]: <seaborn.axisgrid.PairGrid at 0x2403c9fb940>
```

In [467]: `X = data['floors']`
`X.head()`

Out[467]:

0	1.0
1	2.0
2	1.0
3	1.0
4	1.0

Name: floors, dtype: float64

In [399]: `X`

Out[399]:

0	1.0
1	2.0
2	1.0
3	1.0
4	1.0
...	...
21608	3.0
21609	2.0
21610	2.0
21611	2.0
21612	2.0


Name: floors, Length: 21613, dtype: float64

In [468]: `y = data['price']`
`y.head()`

Out[468]:

0	221900.0
1	538000.0
2	180000.0
3	604000.0
4	510000.0

Name: price, dtype: float64



In [441]: y

```
Out[441]: 0      221900.0
          1      538000.0
          2      180000.0
          3      604000.0
          4      510000.0
          ...
        21608    360000.0
        21609    400000.0
        21610    402101.0
        21611    400000.0
        21612    325000.0
        Name: price, Length: 21613, dtype: float64
```

In [575]: `from sklearn.model_selection import train_test_split`

In [576]: `X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_state=42)`

In [577]: `print(type(X_train))`
`print(type(y_train))`
`print(type(X_test))`
`print(type(y_test))`

```
<class 'pandas.core.series.Series'>
<class 'pandas.core.series.Series'>
<class 'pandas.core.series.Series'>
<class 'pandas.core.series.Series'>
```

```
In [578]: print(X_train.shape)
          print(y_train.shape)
          print(X_test.shape)
          print(y_test.shape)
```

```
(14480,)
(14480,)
(7133,)
(7133,)
```

```
In [579]: X_train = X_train[:,np.newaxis]
          X_test = X_test[:,np.newaxis]
```

C:\Users\my pc\AppData\Local\Temp\ipykernel_20152\2733214152.py:1: FutureWarning:

Support for multi-dimensional indexing (e.g. `obj[:, None]`) is deprecated and will be removed in a future version. Convert to a numpy array before indexing instead.

C:\Users\my pc\AppData\Local\Temp\ipykernel_20152\2733214152.py:2: FutureWarning:

Support for multi-dimensional indexing (e.g. `obj[:, None]`) is deprecated and will be removed in a future version. Convert to a numpy array before indexing instead.

```
In [580]: print(X_train.shape)
          print(y_train.shape)
          print(X_test.shape)
          print(y_test.shape)
```

```
(14480, 1)
(14480,)
(7133, 1)
(7133,)
```

```
In [581]: from sklearn.linear_model import LinearRegression
```

```
In [582]: ln = LinearRegression()
```

```
In [583]: ln.fit(X_train,y_train)
```

```
Out[583]: LinearRegression()
```

```
In [584]: ln.intercept_
```

```
Out[584]: 285395.0290051142
```

```
In [585]: print(ln.intercept_)
```

```
285395.0290051142
```

```
In [586]: print(ln.coef_)
```

```
[168371.20436773]
```

```
In [587]: x_pred = ln.predict(X_test)
x_pred
```

```
Out[587]: array([622137.43774057, 453766.23337284, 622137.43774057, ...,
        622137.43774057, 790508.6421083 , 453766.23337284])
```

```
In [588]: import matplotlib.pyplot as plt
%matplotlib inline
y_pred = ln.predict(X_test)
c = [ i for i in range(1,67,1)]
fig = plt.figure()
plt.plot(c, y_test,color = 'blue', lw = 2.5, linestyle = '-')
plt.plot(c, y_pred,color = 'r', lw = 2.5, linestyle = '-')
```

```

-----
ValueError                                Traceback (most recent call last)
~\AppData\Local\Temp\ipykernel_20152\3983697950.py in <module>
      4 c = [ i for i in range(1,67,1)]
      5 fig = plt.figure()
----> 6 plt.plot(c, y_test,color = 'blue', lw = 2.5, linestyle = '-')
      7 plt.plot(c, y_pred,color = 'r', lw = 2.5, linestyle = '-')

~\anaconda3\lib\site-packages\matplotlib\pyplot.py in plot(scalex, scaley, data, *args, **kwargs)
    2767 @_copy_docstring_and_deprecators(Axes.plot)
    2768 def plot(*args, scalex=True, scaley=True, data=None, **kwargs):
-> 2769     return gca().plot(
    2770         *args, scalex=scalex, scaley=scaley,
    2771         **({"data": data} if data is not None else {}), **kwargs)


~\anaconda3\lib\site-packages\matplotlib\axes\_axes.py in plot(self, scalex, scaley, data, *args, **kwargs)
    1630     """
    1631     kwargs = cbook.normalize_kwargs(kwargs, mlines.Line2D)
-> 1632     lines = [*self._get_lines(*args, data=data, **kwargs)]
    1633     for line in lines:
    1634         self.add_line(line)

~\anaconda3\lib\site-packages\matplotlib\axes\_base.py in __call__(self, data, *args, **kwargs)
    310         this += args[0],
    311         args = args[1:]
--> 312         yield from self._plot_args(this, kwargs)
    313
    314     def get_next_color(self):

~\anaconda3\lib\site-packages\matplotlib\axes\_base.py in _plot_args(self, tup, kwargs, return_kwargs)
    496
    497     if x.shape[0] != y.shape[0]:
--> 498         raise ValueError(f"x and y must have same first dimension, but "
    499                         f"have shapes {x.shape} and {y.shape}")
    500     if x.ndim > 2 or y.ndim > 2:

ValueError: x and y must have same first dimension, but have shapes (66,) and (7133,)

```



In []: