

CNN : CONVENTIONAL NEURAL NETWORKS IN TENSOR FLOW :

PAGE 140

In []:

In []:

1. STAGE 1 : INSTALLING DEPENDENCIES AND NOTEBOOK GPU SETUP:

```
In [1]: # "pip install tensorflow" - in ANACONDA POWERSHELL PROMPT :  
# pip install matplotlib-venn(successfully installed)  
# !apt-get -qq install -y libfluidsynth1(it didn't worked)
```

2. STAGE 2 : IMPORTING DEPENDENCIES FOR THE PROJECT :

PAGE 154 BOOK 4

```
In [41]: import tensorflow as tf
import matplotlib.pyplot as plt

from tensorflow.keras.datasets import cifar10

from tensorflow.keras.layers import Conv2D, Dense, Flatten # i personally added this syntax

%matplotlib inline
tf.__version__
```

Out[41]: '2.12.0'

STAGE 3 : DATA PREPROCESSING :

LOADING THE 'cifar10' DATASET :

PAGE 155

SETTING CLASS NAMES FOR THE DATASET :

```
In [42]: class_names=['airplane','automobile','bird','cat','deer','dog','frog','horse','ship','truck']
```

LOADING THE DATASET :

PAGE 155

```
In [43]: (X_train,y_train),(X_test,y_test)=cifar10.load_data()
```

IMAGE NORMALISATION :

PAGE 155 B00K 4

```
In [44]: X_train = X_train/255.0
```


```
In [45]: X_train.shape
```

```
Out[45]: (50000, 32, 32, 3)
```

```
In [46]: X_test = X_test/255.0
```

```
In [47]: plt.imshow(X_test[10])
```

```
Out[47]: <matplotlib.image.AxesImage at 0x25508886880>
```



STAGE 4 : BUILDING A CONVOLUTIONAL NEURAL NETWORK :

DEFINING THE MODEL :

Sequential - 'S' Capital.

```
In [48]: model = tf.keras.models.Sequential()
```

ADDING THE FIRST 'CNN' LAYER:

PAGE 157

#'CNN' layer hyper-parameters: filters : 32

kernel_size : 3

padding : same

activation : relu

input-shape : (32,32,3)

Conv2D - 'C' and 'D' Capitals

```
In [49]: model.add(tf.keras.layers.Conv2D(filters = 32, kernel_size = 3, padding = "same", activation="relu", input_shape=[32,32,3]))
```

ADDING THE SECOND 'CNN' LAYER AND 'MAX POOL' LAYER :

PAGE 157

#'CNN' layer hyper-parameters:

filters : 32

kernel_size : 3

padding : same

activation : relu

Conv2D - 'C' and 'D' Capitals

#'Max Pool' layer hyper-parameters:

pool_size : 2

strides : 2

padding : valid

```
In [50]: model.add(tf.keras.layers.Conv2D(filters=32, kernel_size=3, padding="same", activation="relu"))
```

```
In [51]: model.add(tf.keras.layers.MaxPool2D(pool_size=2, strides=2, padding='valid'))
```

ADDING THE THIRD 'CNN' LAYER:

PAGE 158

#'CNN' layer hyper-parameters:

filters : 64

kernel_size : 3

padding : same

activation : relu

input-shape : (32,32,3)

Conv2D - 'C' and 'D' Capitals

```
In [52]: model.add(tf.keras.layers.Conv2D(filters = 64, kernel_size= 3, padding = "same", activation = "relu"))
```

ADDING THE FOURTH 'CNN' LAYER AND 'MAX POOL' LAYER :

PAGE 158

#'CNN' layer hyper-parameters:

filters : 64

kernel_size : 3

padding : same

activation : relu

Conv2D - 'C' and 'D' Capitals

#'Max Pool' layer hyper-parameters:

pool_size : 2

strides : 2

padding : valid

```
In [53]: model.add(tf.keras.layers.Conv2D(filters=64, kernel_size = 3, padding = "same", activation = "relu"))
```

```
In [54]: model.add(tf.keras.layers.MaxPool2D(pool_size = 2, strides = 2, padding = 'valid'))
```

ADDING THE 'FLATTEN LAYER' :

PAGE 159

```
In [55]: model.add(tf.keras.layers.Flatten())
```

ADDING THE FIRST DENSE LAYER :

#'DENSE' LAYER HYPER-PARAMETERS :

units/neurons : 128

activation : Softmax

```
In [56]: model.add(tf.keras.layers.Dense(units = 128, activation = 'Softmax'))
```

ADDING THE SECOND DENSE LAYER : (OUTPUT LAYER)

#'DENSE' LAYER HYPER-PARAMETERS :

units/neurons : 10 (NUMBER OF CLASSES)

activation : Softmax

```
In [57]: model.add(tf.keras.layers.Dense(units = 10, activation = 'Softmax'))
```



```
In [58]: model.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
=====		
conv2d_2 (Conv2D)	(None, 32, 32, 32)	896
conv2d_3 (Conv2D)	(None, 32, 32, 32)	9248
max_pooling2d_1 (MaxPooling 2D)	(None, 16, 16, 32)	0
conv2d_4 (Conv2D)	(None, 16, 16, 64)	18496
conv2d_5 (Conv2D)	(None, 16, 16, 64)	36928
max_pooling2d_2 (MaxPooling 2D)	(None, 8, 8, 64)	0
flatten_1 (Flatten)	(None, 4096)	0
dense_2 (Dense)	(None, 128)	524416
dense_3 (Dense)	(None, 10)	1290
=====		
Total params: 591,274		
Trainable params: 591,274		
Non-trainable params: 0		
=====		

STAGE 5 : COMPILING THE MODEL :

PAGE 160

#sparse_categorical_accuracy -> checks to see if the maximal true value is equal to the Index of the maximal predicted value.

```
In [59]: model.compile(loss = "sparse_categorical_crossentropy", optimizer = "Adam", metrics = ["sparse_categorical_accuracy"])
```

STAGE 6 : TRAINING THE MODEL :

PAGE 160

```
In [60]: model.fit(X_train, y_train, epochs = 5)
```

```
Epoch 1/5  
1563/1563 [=====] - 137s 86ms/step - loss: 2.2603 - sparse_categorical_accuracy: 0.1442  
Epoch 2/5  
1563/1563 [=====] - 133s 85ms/step - loss: 2.2833 - sparse_categorical_accuracy: 0.1205  
Epoch 3/5  
1563/1563 [=====] - 134s 86ms/step - loss: 2.3029 - sparse_categorical_accuracy: 0.0965  
Epoch 4/5  
1563/1563 [=====] - 133s 85ms/step - loss: 2.3028 - sparse_categorical_accuracy: 0.0989  
Epoch 5/5  
1563/1563 [=====] - 134s 86ms/step - loss: 2.3028 - sparse_categorical_accuracy: 0.0987
```

```
Out[60]: <keras.callbacks.History at 0x25508a11580>
```

STAGE 7 : MODEL EVALUATION AND PREDICTION :

PAGE 161

```
In [61]: test_loss, test_accuracy = model.evaluate(X_test, y_test)
```

```
313/313 [=====] - 8s 25ms/step - loss: 2.3029 - sparse_categorical_accuracy: 0.1000
```

```
In [62]: # Here we are identifying the Accurate test :  
  
print("Test Accuracy : {}".format(test_accuracy))
```

Test Accuracy : 0.10000000149011612

In []:

In []:

HAND WRITTING DIGIT RECOGNIZER USING 'DEEP LEARNING' :

PAGE 162

```
In [63]: # In this Project, we will try to identify Hand Writing digits by using the power of 'Deep Learning'.  
# We will be Working on the 'MNIST' dataset to create a DEEP LEARNING CLASSIFICATION MODEL & See how our Model perform  
# in accurately Acurately Predicting Images with the Correct digit notation.  
  
# 'MNIST' is a database of Hand Writing Digits made up of a Training Set of 60000 examples and a Testset of 10000 exa  
# The Training examples are annotated by Humans wih the Correct Answer.  
# For Instance, if the Handwritting digit is the number '3', then'3 is simply the Label associated with that example.
```

```
In [64]: # We will Train the model with the samples available in the training set and then use the test set to evaluate -  
# how well our neural network has Learned to recognise digits.  
  
# Now Let's create a Convolution Neural Network to Solve this Problem.
```

```
In [65]: # The 'MNIST database' (Modified National Institute of Standards and Technology database)  
# -> is a Large database of handwritten digits that is commonly used for training various image processing systems.
```

1. 'import Libraries' :

page 163

```
In [66]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [67]: import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPool2D, Flatten, Dense, Dropout
from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras.utils import to_categorical
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score, precision_score, recall_score
%matplotlib inline
```

```
In [68]: # We can use the following code snippet to check which version of TensorFlow is installed in our system:
# This will print the version of TensorFlow that is currently installed on our system.
# For example, if we are running version 2.4.1 of TensorFlow, the below code will output "2.4.1".
```

```
tf.__version__
```

```
Out[68]: '2.12.0'
```

```
In [ ]:
```

2. 'Import MNIST Dataset' :

page 164

```
In [90]: from tensorflow.keras.datasets import mnist
```

In [70]: *# In this one also we have an inbuilt dataset, where the dataset shapes are 'X' 60000 of 10000 and 'y' 60000 of 10000*

```
(X_train,y_train),(X_test,y_test) = mnist.load_data()  
print("shape of X_train:",X_train.shape)  
print("shape of X_test:",X_test.shape)  
print("shape of y_train:",y_train.shape)  
print("shape of y_test:",y_test.shape)
```


```
shape of X_train: (60000, 28, 28)  
shape of X_test: (10000, 28, 28)  
shape of y_train: (60000,)  
shape of y_test: (10000,)
```

3. Each MNIST Image is in GrayScale and consists of 28*28 pixels :

page 164

```
In [71]: from PIL import Image  
Image.open('mnist.png')
```

Out[71]:




4. Load Sample Image from the 'MNIST' Dataset :

page 165

In [72]: `plt.figure(figsize=(20,20))`

```
plt.subplot(1,10,1)
plt.imshow(X_train[0])
plt.subplot(1,10,2)
plt.imshow(X_train[1])
plt.subplot(1,10,3)
plt.imshow(X_train[2])
plt.subplot(1,10,4)
plt.imshow(X_train[3])
plt.subplot(1,10,5)
plt.imshow(X_train[4])
plt.subplot(1,10,6)
plt.imshow(X_train[5])
plt.subplot(1,10,7)
plt.imshow(X_train[6])
plt.subplot(1,10,8)
plt.imshow(X_train[7])
plt.subplot(1,10,9)
plt.imshow(X_train[8])
plt.subplot(1,10,10)
plt.imshow(X_train[9])
```

Out[72]: `<matplotlib.image.AxesImage at 0x2550a00c3d0>`





```
In [73]: # Here, This particular data, in the image pattern, we are changing in to an a pixel.  
  
# The image data cannot be fed directly into the model.  
# So, we need to perform some operations and process the data to make it ready for our Neural Networks.  
# The Dimensions of the Training data is (60000,28,28).  
# The CNN Model will require one more dimension so we reshape the matrix to shape (60000,28,28,1).  
# This extra dimension is for the Color Channel for Grayscale images like MNIST, it's value is 1.  
# For Color images, the Channel Value is '3' Corresponding to 'Red, Green & Blue(RGB)'.
```

```
In [74]: plt.figure(figsize=(20,20))

plt.subplot(1,10,1)
plt.imshow(X_train[150])
plt.subplot(1,10,2)
plt.imshow(X_train[162])
plt.subplot(1,10,3)
plt.imshow(X_train[178])
plt.subplot(1,10,4)
plt.imshow(X_train[193])
plt.subplot(1,10,5)
plt.imshow(X_train[205])
plt.subplot(1,10,6)
plt.imshow(X_train[3978])
plt.subplot(1,10,7)
plt.imshow(X_train[456])
plt.subplot(1,10,8)
plt.imshow(X_train[7896])
plt.subplot(1,10,9)
plt.imshow(X_train[57])
plt.subplot(1,10,10)
plt.imshow(X_train[31897])
```

Out[74]: <matplotlib.image.AxesImage at 0x2550e4363d0>



5. DATA PREPROCESSING :

PAGE 166

```
In [75]: print("Shape X_train: ", X_train.shape)
         print("Shape X_test: ", X_test.shape)
```

```
Shape X_train: (60000, 28, 28)
Shape X_test: (10000, 28, 28)
```

```
In [76]: X_train = X_train.reshape(X_train.shape[0],28,28,1)
         X_test = X_test.reshape(X_test.shape[0],28,28,1)
         input_shape = (28,28,1)
         print("Shape of X_train: ", X_train.shape)
         print("Shape of X_test: ", X_test.shape)
```

```
Shape of X_train: (60000, 28, 28, 1)
Shape of X_test: (10000, 28, 28, 1)
```

6. ONE HOT ENCODING OF TARGET LABELS :

PAGE 167

CATEGORICAL DATA : (NON-NUMERICAL)

```
In [77]: # We are going to use 'OHE(NON-NUMERICAL)' as a Sample Tool to encode information used inside Neural Networks.
         # In many applications it is Convenient to Transform Categorical(non-numerical) features into 'Numerical Variables'.
         # For Instance, The Categorical Feature 'digit' with Value 'd' in [0 to 9] can be encoded into 'Binary Vector' with 10
         # Which always has '0' Value, except the 'd'th Position, Where a '1' is present.

         # For Example, the digit '3' can be encoded as [0,0,0,1,0,0,0,0,0,0],
         # This type of Representation is called 'ONE-HOT ENCODING(OHE)', OR Sometimes Simply ONE-HOT, and is very Common in de
         # -mining when the Learning Algorithm is Specialized in dealing with 'Numerical Functions'.
```

```
In [78]: y_train[0:11]
```

```
Out[78]: array([5, 0, 4, 1, 9, 2, 1, 3, 1, 4, 3], dtype=uint8)
```

```
In [79]: y_cat_train = to_categorical(y_train,10)
y_cat_test = to_categorical(y_test,10)
y_cat_train[0:11]
```

```
Out[79]: array([[0., 0., 0., 0., 0., 1., 0., 0., 0., 0.],
 [1., 0., 0., 0., 0., 0., 0., 0., 0., 0.],
 [0., 0., 0., 0., 1., 0., 0., 0., 0., 0.],
 [0., 1., 0., 0., 0., 0., 0., 0., 0., 0.],
 [0., 0., 0., 0., 0., 0., 0., 0., 0., 1.],
 [0., 0., 1., 0., 0., 0., 0., 0., 0., 0.],
 [0., 1., 0., 0., 0., 0., 0., 0., 0., 0.],
 [0., 0., 0., 1., 0., 0., 0., 0., 0., 0.],
 [0., 1., 0., 0., 0., 0., 0., 0., 0., 0.],
 [0., 0., 0., 0., 1., 0., 0., 0., 0., 0.],
 [0., 0., 0., 1., 0., 0., 0., 0., 0., 0.]], dtype=float32)
```

7. SCALING FEATURE DATA :

PAGE 169

```
In [80]: # NEURAL NETWORKS Works Well, When the Feature Values lie between '0 to 1'.

# Hence, We will Scale the dataset by Simply Dividing Each Value by 255.

# The Value for Each Pixel is case of 'Gray Scale Images range from '0(white)' to '255(Black)'.
```

[illegible]

```
X_train = X_train.astype('float32')
X_test = X_test.astype('float32')

X_train /= 255
X_test /= 255
```

In [83]: X_train[0]

```
[0.          ],
[0.          ],
[0.          ],
[0.          ],
[0.          ],
[0.          ],
[0.3137255 ],
[0.6117647 ],
[0.41960785],
[0.99215686],
[0.99215686],
[0.8039216 ],
[0.04313726],
[0.          ],
[0.16862746],
[0.6039216 ],
[0.          ],
[0.          ],
[0.          ],
[0.          ],
[0.          ],
```

8. MODEL CREATION :

PAGE 170

```
In [84]: # Now we will Create our CNN Model.
# A CNN Model generally consists of Convolutional and Pooling Layers.
# It Works better for Data that are represented as Grid Structures, this is the reason why CNN Works well,
# for Image Classification Problems.
```

```
In [95]: model = Sequential()

model.add(Conv2D(32, kernel_size = (3,3), activation = 'relu', input_shape = input_shape))
model.add(Conv2D(64, (3,3), activation = 'relu' ))
model.add(MaxPool2D(pool_size = (2,2)))
model.add(Dropout(0.25))
model.add(Flatten())

model.add(Dense(256, activation = 'relu'))
model.add(Dropout(0.25))
model.add(Dense(10, activation = 'softmax'))
model.compile(loss = 'categorical_crossentropy', optimizer = 'adam', metrics = ['accuracy'])
```

9. Model Summary :

page 171

```
In [96]: model.summary()
```

Model: "sequential_5"

Layer (type)	Output Shape	Param #
=====		
conv2d_8 (Conv2D)	(None, 26, 26, 32)	320
conv2d_9 (Conv2D)	(None, 24, 24, 64)	18496
max_pooling2d_5 (MaxPooling 2D)	(None, 12, 12, 64)	0
dropout_6 (Dropout)	(None, 12, 12, 64)	0
flatten_4 (Flatten)	(None, 9216)	0
dense_8 (Dense)	(None, 256)	2359552
dropout_7 (Dropout)	(None, 256)	0
dense_9 (Dense)	(None, 10)	2570
=====		
Total params: 2,380,938		
Trainable params: 2,380,938		
Non-trainable params: 0		

10. ADDING 'EARLY STOPPING' :

PAGE 171

```
In [97]: early_stop = EarlyStopping(monitor = 'val_loss', patience = 2)
```

11. MODEL TRAINING :

PAGE 172

In [98]: *# And if we observe cross check, it looks like, Model training what we have seen in the 'ANN'.*

```
In [99]: model.fit(X_train, y_cat_train, epochs = 50, callbacks = [early_stop], validation_data = (X_test, y_cat_test))

print("The Model has Successfully Trained ")
model.save('mnist.h5')
print("Saving the Model as mnist.h5 ")
```

Epoch 1/50

1875/1875 [=====] - 102s 54ms/step - loss: 0.1270 - accuracy: 0.9601 - val_loss: 0.0402 - val_accuracy: 0.9869

Epoch 2/50

1875/1875 [=====] - 106s 57ms/step - loss: 0.0476 - accuracy: 0.9850 - val_loss: 0.0349 - val_accuracy: 0.9889

Epoch 3/50

1875/1875 [=====] - 112s 60ms/step - loss: 0.0325 - accuracy: 0.9898 - val_loss: 0.0296 - val_accuracy: 0.9889

Epoch 4/50

1875/1875 [=====] - 107s 57ms/step - loss: 0.0246 - accuracy: 0.9921 - val_loss: 0.0256 - val_accuracy: 0.9922

Epoch 5/50

1875/1875 [=====] - 122s 65ms/step - loss: 0.0186 - accuracy: 0.9935 - val_loss: 0.0399 - val_accuracy: 0.9887

Epoch 6/50

1875/1875 [=====] - 139s 74ms/step - loss: 0.0170 - accuracy: 0.9944 - val_loss: 0.0323 - val_accuracy: 0.9908

The Model has Successfully Trained

Saving the Model as mnist.h5

12. MODEL PERFORMANCE DURING TRAINING & VALIDATION :

PAGE 172

```
In [104]: training_metrics = pd.DataFrame(model.history.history)
          training_metrics.columns
```

```
Out[104]: Index(['loss', 'accuracy', 'val_loss', 'val_accuracy'], dtype='object')
```


```
In [105]: training_metrics.head()
```

```
Out[105]:
```

	loss	accuracy	val_loss	val_accuracy
0	0.127006	0.960133	0.040150	0.9869
1	0.047592	0.985017	0.034949	0.9889
2	0.032518	0.989833	0.029571	0.9889
3	0.024573	0.992100	0.025625	0.9922
4	0.018638	0.993533	0.039902	0.9887


```
In [106]: training_metrics[['loss', 'val_loss']].plot()
```

Out[106]: <AxesSubplot:>



```
In [107]: training_metrics[['accuracy', 'val_accuracy']].plot()
```

Out[107]: <AxesSubplot:>



NOW ACCURACY TEST :

```
In [108]: score = model.evaluate(X_test, y_cat_test, verbose = 0)
print('Test loss : ', score [0])
print('Test accuracy : ', score [1])
```

Test loss : 0.03234346956014633
Test accuracy : 0.9908000230789185

13. MODEL PREDICTIONS :

PAGE 174

```
In [113]: predictions = np.argmax(model.predict(X_test), axis = -1)
          print(classification_report(y_test, predictions))
```

313/313 [=====] - 4s 13ms/step

	precision	recall	f1-score	support
0	0.99	1.00	0.99	980
1	1.00	0.99	1.00	1135
2	1.00	0.99	0.99	1032
3	0.98	1.00	0.99	1010
4	0.99	0.98	0.99	982
5	0.99	0.99	0.99	892
6	0.99	0.98	0.99	958
7	0.99	1.00	0.99	1028
8	0.99	0.99	0.99	974
9	0.98	0.99	0.99	1009
accuracy			0.99	10000
macro avg	0.99	0.99	0.99	10000
weighted avg	0.99	0.99	0.99	10000

```
In [115]: print(confusion_matrix(y_test,predictions))
```


```
[[ 977    0    0    1    0    0    2    0    0    0]
 [   1 1127    1    3    0    0    1    1    1    0]
 [   1    1 1023    1    0    0    0    3    3    0]
 [   0    0    0 1009    0    1    0    0    0    0]
 [   0    0    0    0 966    0    4    0    1   11]
 [   0    0    0    9    0 879    1    0    1    2]
 [   4    2    0    1    2    4 943    0    2    0]
 [   0    0    2    1    0    0    0 1024    1    0]
 [   1    0    1    5    0    0    0    1 963    3]
 [   0    0    0    2    4    0    0    3    3 997]]
```

14. VISUALIZING CONFUSION MATRIX :

PAGE 175

```
In [118]: plt.figure(figsize = (15,7))
sns.heatmap(confusion_matrix(y_test, predictions),annot=True)
```

Out[118]: <AxesSubplot:>




15. PREDICTING INDIVIDUAL IMAGES :

PAGE 175

```
In [120]: # Here, This is predicting individually.  
# It is showing the Image pattern.
```

```
new_img = X_test[95]  
plt.imshow(new_img)
```

```
Out[120]: <matplotlib.image.AxesImage at 0x25509f71d90>
```



```
In [121]: y_test[95]
```

```
Out[121]: 4
```




```
In [123]: np.argmax(model.predict(new_img.reshape(1,28,28,1)), axis = 1)
```

```
1/1 [=====] - 0s 20ms/step
```

```
Out[123]: array([4], dtype=int64)
```

```
In [124]: new_img2 = X_test[0]  
plt.imshow(new_img2)
```

```
Out[124]: <matplotlib.image.AxesImage at 0x2550e529f10>
```



```
In [125]: y_test[0]
```

```
Out[125]: 7
```


```
In [128]: np.argmax(model.predict(new_img2.reshape(1,28,28,1)), axis = -1)
```

```
1/1 [=====] - 0s 29ms/step
```

```
Out[128]: array([7], dtype=int64)
```

```
In [129]: new_img3 = X_test[397]  
plt.imshow(new_img3)
```

```
Out[129]: <matplotlib.image.AxesImage at 0x2550e57c160>
```



```
In [130]: np.argmax(model.predict(new_img3.reshape(1,28,28,1)), axis = 1)
```

```
1/1 [=====] - 0s 40ms/step
```

```
Out[130]: array([5], dtype=int64)
```

```
In [ ]:
```